

REDES NEURONALES CONVOLUTIVAS




MINERÍA DE DATOS

4º Curso. Grado en Ingeniería Informática
5º Curso. Doble Grado Informática/Estadística (INDAT)

Tratamiento numérico de imágenes

MNIST

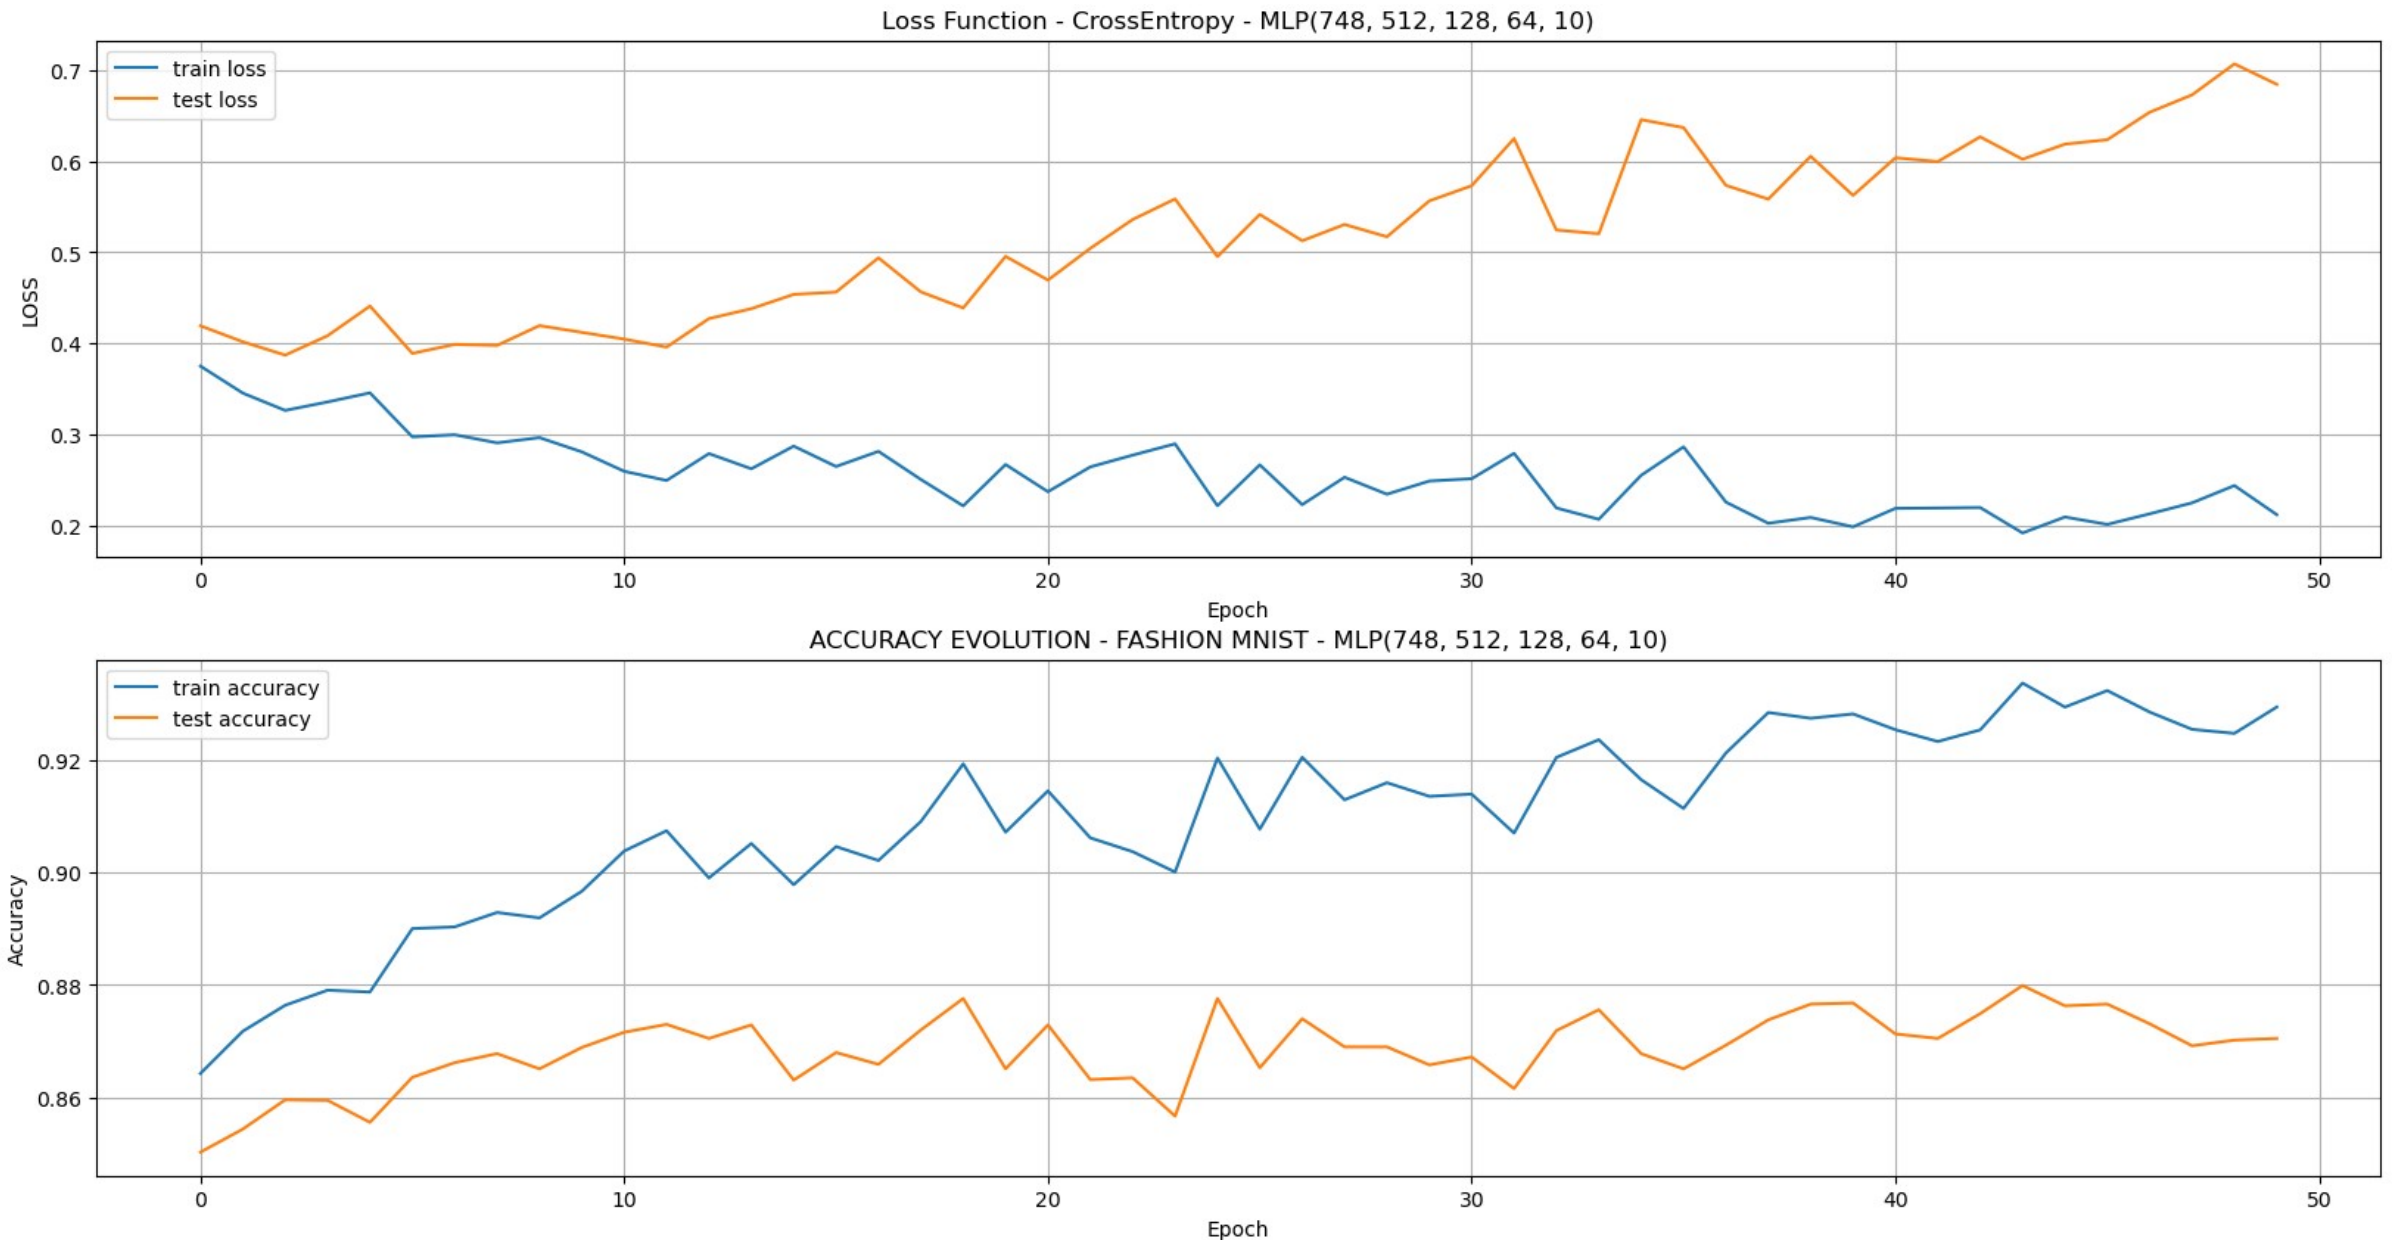


- Ineficiencia al tener que convertir tipos **int→float**:
Float32 (por defecto en Deep Learning) y Punto Fijo (no flotante)
- Pérdida de información al **aplanar**

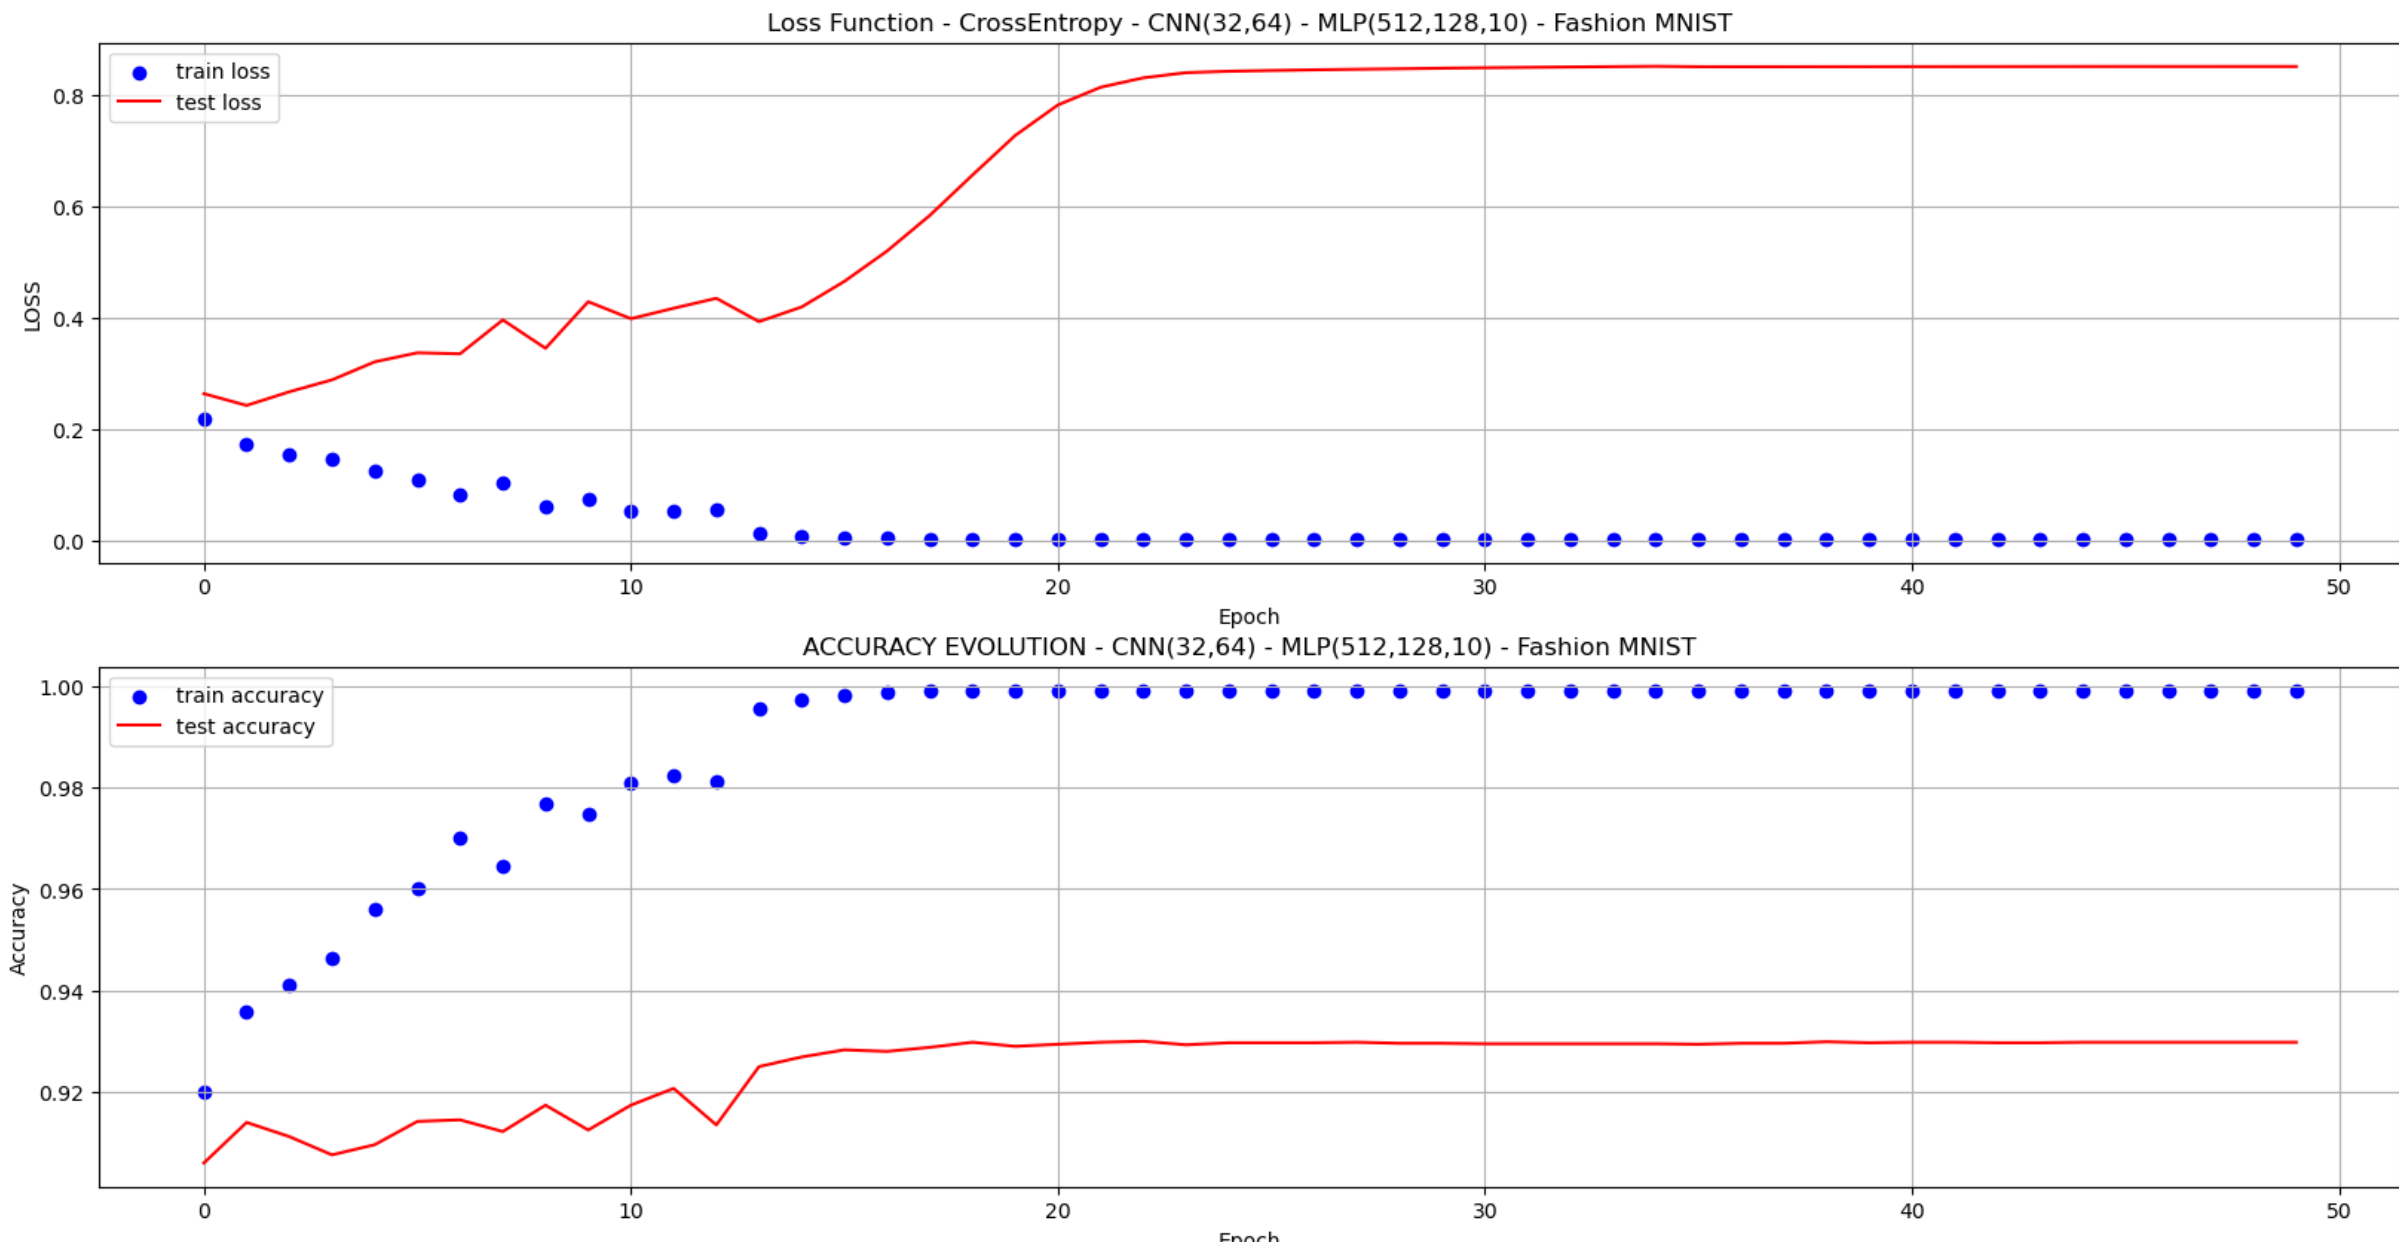
Redes convolutivas:

- Extraen características 2D
- Para ello aplican un núcleo, que es una matriz 2D

Fashion MNIST: MLP 512-128-64-10




CNN - VGG



Imágenes a color


$$I_R(m, n, 1) = \begin{bmatrix} r_{11} & r_{12} & \dots & r_{1n} \\ r_{21} & r_{22} & \dots & r_{2n} \\ \vdots & \vdots & \ddots & \vdots \\ r_{m1} & r_{m2} & \dots & r_{mn} \end{bmatrix} \quad I_G(m, n, 2) = \begin{bmatrix} g_{11} & g_{12} & \dots & g_{1n} \\ g_{21} & g_{22} & \dots & g_{2n} \\ \vdots & \vdots & \ddots & \vdots \\ g_{m1} & g_{m2} & \dots & g_{mn} \end{bmatrix} \quad I_B(m, n, 3) = \begin{bmatrix} b_{11} & b_{12} & \dots & b_{1n} \\ b_{21} & b_{22} & \dots & b_{2n} \\ \vdots & \vdots & \ddots & \vdots \\ b_{m1} & b_{m2} & \dots & b_{mn} \end{bmatrix}$$

1=R
2=G
3=B



- Vector tensor 2D: (samples, features)
- Series temporales tensor 3D: (samples, timesteps, features)
- **Imágenes tensors 4D, ejemplos:**
(sample, height, width, channel) or (sample, channel, height, width)
- Vídeo tensor 5D:
(sample, frame, height, width, channel) o (sample, frame, channel, height, width)


Convolución 2D:



AIGeekProgrammer.com © 2019


Se aprovecha la información 2D

Convolución 2D (RGB)




Capa convolucional – entrada: 3 canales – salida: 2 canales

Avance o “Stride”





- Núcleo 3x3
- Imagen original arriba
- Padding
- Stride = 1



- Núcleo 5x5
- Imagen original abajo
- No Padding
- Stride = 4

Avance del recorrido del núcleo



- Tamaño: 3x3, 5x5, 7x7, 9x9, etc.
- Avance (**stride**):
 - De uno en uno (o más)
 - Hasta del tamaño del núcleo
- Se recorre izquierda a derecha
- De arriba a abajo
- Avance horizontal/vertical indep.
- Al llegar a los extremos:
 - Si no hay pixels suficientes
 - Se rellenan artificialmente (**padding**) con pixeles transparentes.

Relleno o “Padding”

0	0	0	0	0	0	0
0	60	113	56	139	85	0
0	73	121	54	84	128	0
0	131	99	70	129	127	0
0	80	57	115	69	134	0
0	104	126	123	95	130	0
0	0	0	0	0	0	0

Kernel

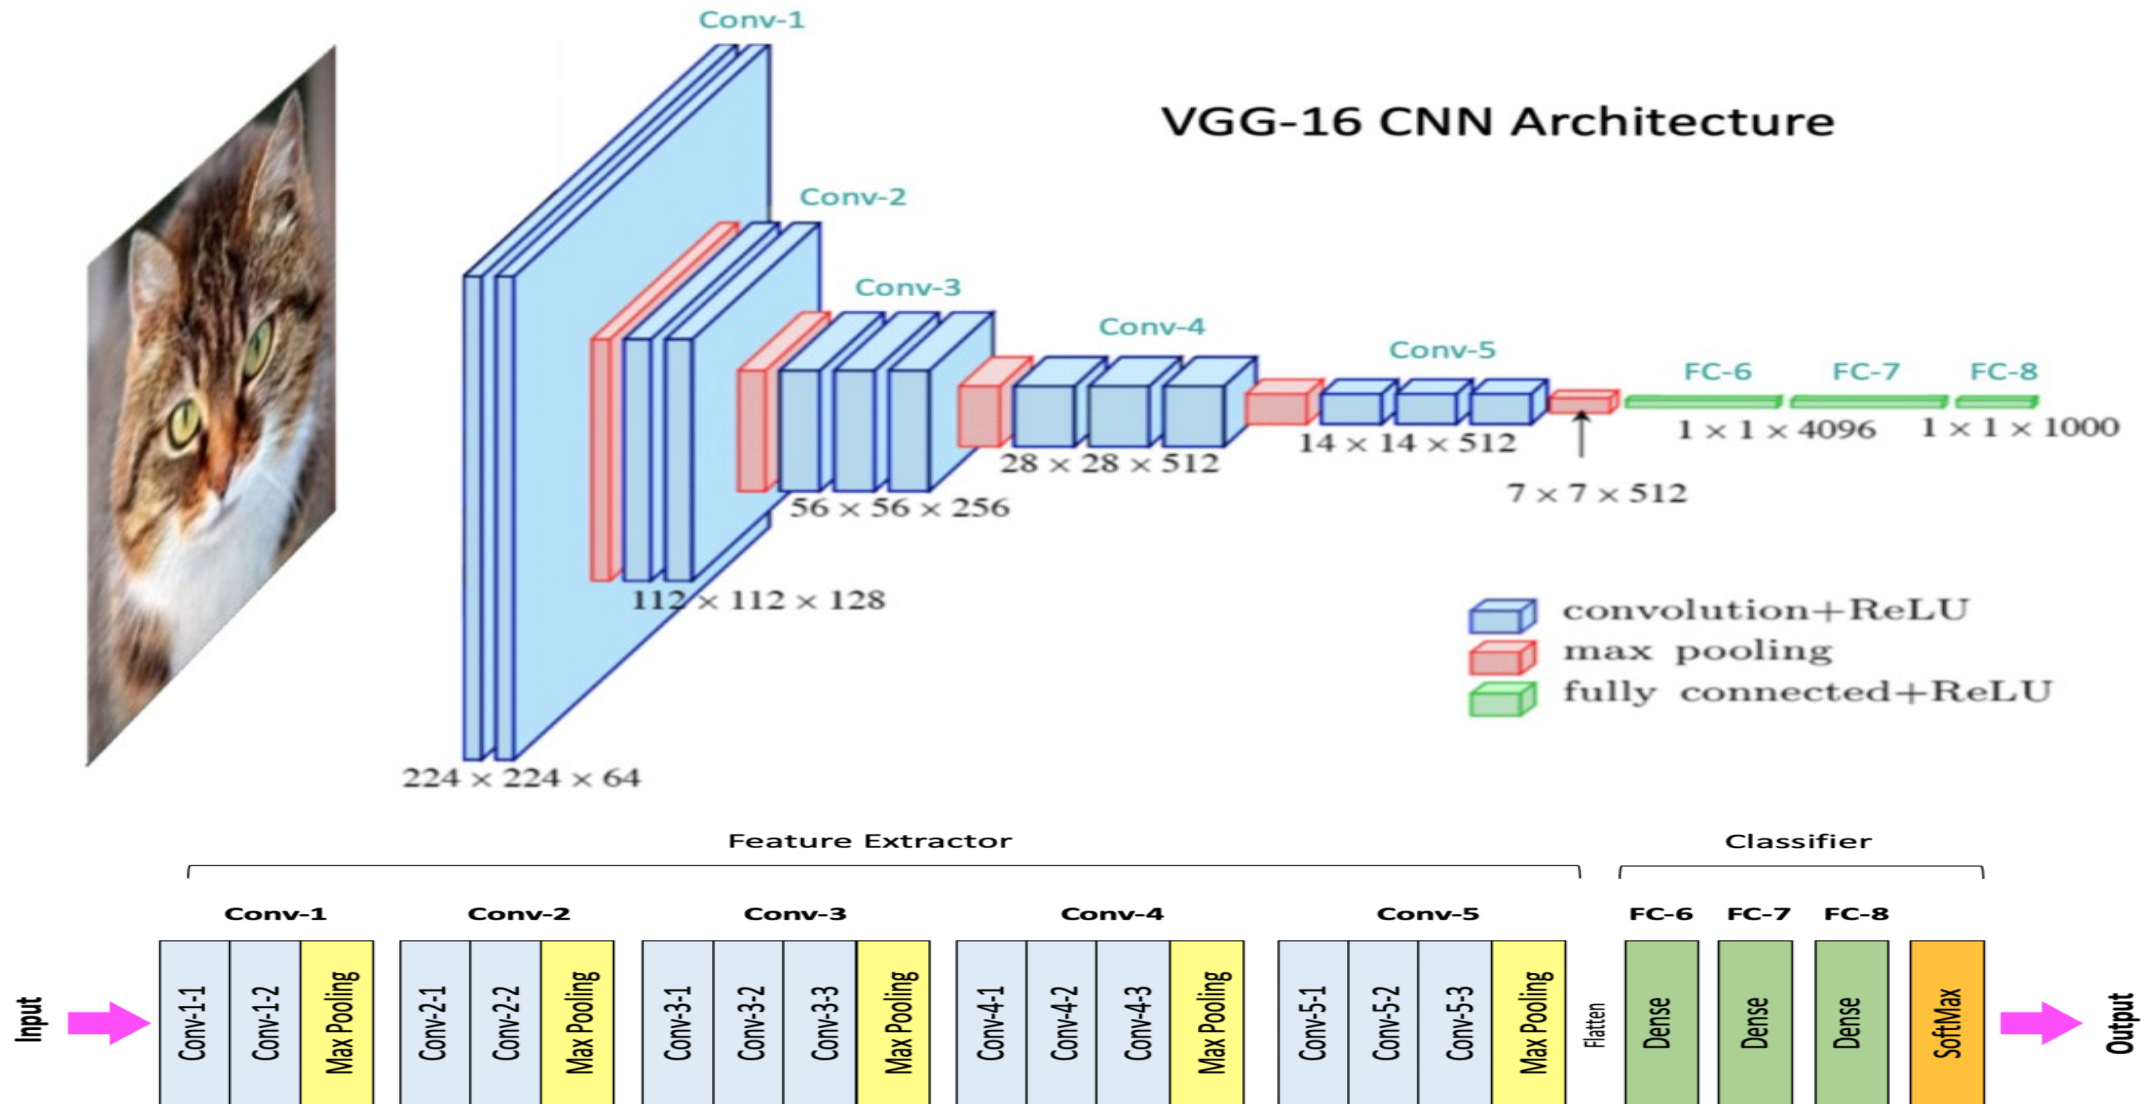
0	-1	0
-1	5	-1
0	-1	0

114				

Relleno más habitual: un borde de tamaño ajustado de ceros ajustado al avance

Rellenos más complejos: basados en interpolaciones

Ejemplo CNN: VGG-16 (channel->last)




Pooling

Max Pooling


Take the **highest** value from the area covered by the kernel

Example: Kernel of size 2 x 2; stride=(2,2)





Average Pooling

Calculate the **average** value from the area covered by the kernel




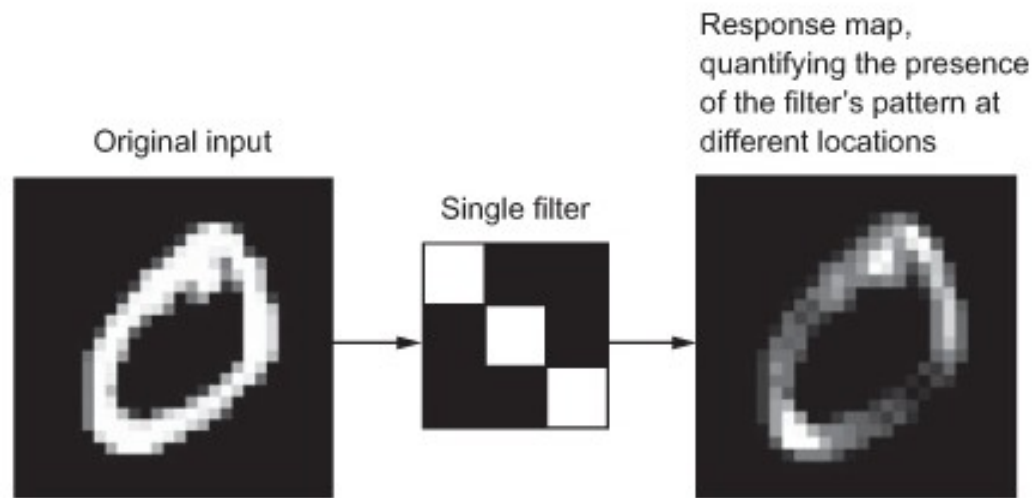
- MaxPooling es el más usado
- El tamaño del núcleo habitual es 2x2
 - Podría ser 3x4, por ejemplo, pero no aparece en las implementaciones
- El avance normal igual a tamaño del núcleo (NxN)
 - Nunca supera a N
 - El más pequeño es uno

Motivación de la Convolución 2D



- Reconocimiento de formas simples en las capas más profundas
- Composición de éstas en la sucesivas capas hasta formar patrones (reconocer) más complejos

Tratamiento matricial de la Convolución



Datasets de Pytorch y Dataloaders

```
train_set = torchvision.datasets.FashionMNIST(root = './data/FashionMNIST', download = True, train = True,  
                                              transform = transforms.Compose([transforms.ToTensor()]))  
  
test_set = torchvision.datasets.FashionMNIST(root = './data/FashionMNIST', download=True, train=False,  
                                              transform = transforms.Compose([transforms.ToTensor()]))
```

<https://pytorch.org/vision/0.9/transforms.html>

```
train_loader = torch.utils.data.DataLoader(train_set, shuffle=True, batch_size=bandeja)  
train_accuracy_loader = torch.utils.data.DataLoader(train_set, batch_size=60000)  
test_loader = torch.utils.data.DataLoader(test_set, batch_size=10000)
```

Datasets y Dataloaders en pytorch

Origen de los datos:

- Dataframes
- Matrices (numpy)
- Tensores (torch)
- Ficheros



Dataset:

- Empareja entrada/salida
- Operaciones



Dataloaders:

Suministro: Memoria o disco
Organiza en bandejas

Datasets creados a partir de tensores

```
class Flatten_Dataset(Dataset):
    def __init__(self, images, labels):
        self.images = images
        self.labels = labels
    def __len__(self):
        return self.images.shape[0]
    def __getitem__(self, idx):
        return self.images[idx], self.labels[idx]
```

```
train_images = torch.zeros((len(train_image_set), 28*28), dtype=torch.float32)
train_labels = torch.empty(len(train_image_set), dtype=torch.long)

for i, muestra in enumerate(train_image_set):
    train_images[i] = torch.flatten(muestra[0])
    train_labels[i] = muestra[1]
```

Creación de la Red Neuronal

```
class Network(nn.Module):
    def __init__(self):
        super(Network,self).__init__()
        self.conv1 = nn.Conv2d(in_channels=1,out_channels=16,kernel_size=3)
        self.conv2 = nn.Conv2d(in_channels=16,out_channels=32,kernel_size=3)

        self.conv3 = nn.Conv2d(in_channels=32,out_channels=64,kernel_size=3)
        self.conv4 = nn.Conv2d(in_channels=64,out_channels=128,kernel_size=3)

        self.relu = nn.LeakyReLU()

        self.pool = nn.MaxPool2d(kernel_size=2,stride=2)

        self.fc1 = nn.Linear(in_features=4*4*128,out_features=128)
        self.out = nn.Linear(in_features=128,out_features=10)

    def forward(self,x):
        #first CNN layer
        x = self.conv1(x)
        x = self.relu(x)
        x = self.conv2(x)
        x = self.relu(x)
        x = self.pool(x)

        #second CNN layer
        x = self.conv3(x)
        x = self.relu(x)
        x = self.conv4(x)
        x = self.relu(x)
        x = self.pool(x)

        #mlp hidden layer
        x = torch.flatten(x, start_dim=1)
        x = self.fc1(x)
        x = self.relu(x)

        #output layer
        x = self.out(x)
        return x
```

```
xx, yy = next(iter(train_loader))
```

```
x1 = model.pool(model.relu(model.conv2(model.relu(model.conv1(xx)))))
x2 = model.pool(model.relu(model.conv4(model.relu(model.conv3(x1)))))
x2.shape
```

```
torch.Size([32, 128, 4, 4])
```

```
summary(model, input_size= (1,28,28))
```

Layer (type)	Output Shape	Param #
LeakyReLU-1	[-1, 1, 28, 28]	0
LeakyReLU-2	[-1, 1, 28, 28]	0
Conv2d-3	[-1, 16, 26, 26]	160
LeakyReLU-4	[-1, 16, 26, 26]	0
Conv2d-5	[-1, 32, 24, 24]	4,640
LeakyReLU-6	[-1, 32, 24, 24]	0
MaxPool2d-7	[-1, 32, 12, 12]	0
Conv2d-8	[-1, 64, 10, 10]	18,496
LeakyReLU-9	[-1, 64, 10, 10]	0
Conv2d-10	[-1, 128, 8, 8]	73,856
LeakyReLU-11	[-1, 128, 8, 8]	0
MaxPool2d-12	[-1, 128, 4, 4]	0
Linear-13	[-1, 128]	262,272
LeakyReLU-14	[-1, 128]	0
LeakyReLU-15	[-1, 128]	0
Linear-16	[-1, 10]	1,290

```
Total params: 360,714
```

```
Trainable params: 360,714
```

```
Non-trainable params: 0
```

```
Input size (MB): 0.00
```

```
Forward/backward pass size (MB): 0.73
```

```
Params size (MB): 1.38
```

```
Estimated Total Size (MB): 2.11
```

Aprendizaje por épocas

```
X_train, y_train = next(iter(train_accuracy_loader))
X_test, y_test = next(iter(test_loader))
```

```
for i in range(epochs):
    print("Epoch: ", i)
    model.train(True)
    for images, targets in tqdm(train_loader):
        #making predictions
        y_pred = model(images)

        #calculating loss
        loss = criterion(y_pred, targets.long())

        #backprop
        optimizer.zero_grad()
        loss.backward()
        optimizer.step()

    with torch.no_grad():
        y_train_pred = model(X_train)
        train_accuracy.append(np.mean((y_train == y_train_pred.argmax(dim=1)).numpy()))
        train_loss.append(criterion(y_train_pred, y_train.long()).numpy())
    print('Train Loss: {0:.5f}'.format(train_loss[-1]), '\tTrain Accuracy: {0:.5f}'.format(train_accuracy[-1]))

    with torch.no_grad():
        y_test_pred = (model(X_test))
        test_accuracy.append(np.mean((y_test == y_test_pred.argmax(dim=1)).numpy()))
        test_loss.append(criterion(y_test_pred, y_test.long()).numpy())

    learning_rate.append(optimizer.param_groups[0]['lr'])
    print('Val Loss: {0:.5f}'.format(test_loss[-1]), '\t Test Accuracy: {0:.5f}'.format(test_accuracy[-1]), "\tLearning Rate:", learning_rate[-1])
```

Falta de RAM tasa aciertos

```
with torch.no_grad():
    aciertos = 0
    muestras = 0
    perdidas = 0.0
    for X_train, y_train in train_accuracy_loader:
        y_train_pred = model(X_train)
        aciertos += np.sum((y_train == y_train_pred.argmax(dim=1)).numpy())
        muestras += X_train.shape[0]
        perdidas += criterion(y_train_pred, y_train.long()).numpy()

train_accuracy.append(aciertos/muestras)
train_loss.append(perdidas)

print('Train Loss: {0:.5f}'.format(train_loss[-1]), '\tTrain Accuracy: {0:.5f}'.format(train_accuracy[-1]))

with torch.no_grad():
    aciertos = 0
    muestras = 0
    perdidas = 0.0
    for X_test, y_test in test_loader:
        y_test_pred = (model(X_test))
        aciertos += np.sum((y_test == y_test_pred.argmax(dim=1)).numpy())
        muestras += X_test.shape[0]
        perdidas += criterion(y_test_pred, y_test.long()).numpy()

test_accuracy.append(aciertos/muestras)
test_loss.append(perdidas)

learning_rate.append(optimizer.param_groups[0]['lr'])
print('Val Loss: {0:.5f}'.format(test_loss[-1]), '\t Test Accuracy: {0:.5f}'.format(test_accuracy[-1]), "\tLearning Rate:", learning_rate[-1])
```


```
train_loader = torch.utils.data.DataLoader(train_set, shuffle=True, batch_size=bandeja)
train_accuracy_loader = torch.utils.data.DataLoader(train_set, batch_size=lote_train)
test_loader = torch.utils.data.DataLoader(test_set, batch_size=lote_test)
```


Ejercicio (entrega):


- Clasificación de imágenes:
 - Dataset: MNIST FASHION incluida en PyTorch
 - Arquitectura VGG:
 - CNN(1,16) + CNN(16,32) + MaxPool(2)
 - CNN(32,64) + CNN(64,128) + MaxPool(2)
 - MLP(¿?, 128, 10)
 - Tamaño de bandeja: 32
 - Función de pérdida: Entropía Cruzada
 - Optimizador: Adagrad (tasa de aprendizaje adaptativo)
 - Épocas:20

Resultados Práctica: Fashion-MNIST

Loss Function - CrossEntropy - CNN(16,32,64,128) - MLP(128,10) - Fashion MNIST



ACCURACY EVOLUTION - CNN(16,32,64,128) - MLP(128,10) - Fashion MNIST



VGG Pre-entrenado

- Las imágenes tienen que ser:
224 x 224 x 3 canales (channel last)
- Deben escalarse para que:
 - Media por canal: 0.485, 0.456, 0.406
 - Desviación estándar por canal: 0.229, 0.224, 0.225
- El modelo tiene 14 ~ 15 millones de pesos.
- De todo esto se concluye que el dataset de imágenes no se va a poder suministrar en matrices, como MNIST, CIFAR, etc.

Conjunto de Imágenes en ficheros

- Se organizan en carpetas: “training”, “test” y, a veces, “validate”.
- Dentro de estos **directorios** aparecen otros tantos correspondiente a las **CLASES**
- Ejemplo de clasificación binaria: perros y gatos



dog.4001.jpg



dog.4002.jpg



dog.4003.jpg



dog.4004.jpg



dog.4021.jpg



dog.4022.jpg



dog.4023.jpg



dog.4024.jpg



dog.4041.jpg



dog.4042.jpg



dog.4043.jpg



dog.4044.jpg




cat.4001.jpg



cat.4002.jpg



cat.4003.jpg



cat.4004.jpg



cat.4021.jpg



cat.4022.jpg



cat.4023.jpg



cat.4024.jpg



cat.4041.jpg



cat.4042.jpg



cat.4043.jpg



cat.4044.jpg

Dataset de ficheros

```
train_data_dir = './data/cats-and-dogs/training_set/training_set/'  
test_data_dir = './data/cats-and-dogs/test_set/test_set/'
```

```
class CatsDogs(Dataset):  
    def __init__(self, folder):  
        cats = glob(folder+'/cats/*.jpg')  
        dogs = glob(folder+'/dogs/*.jpg')  
        self.fpaths = cats[:500] + dogs[:500]  
        self.normalize = transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225])  
        from random import shuffle, seed; seed(10);  
        shuffle(self.fpaths)  
        self.targets = [fpath.split('/')[-1].startswith('dog') for fpath in self.fpaths]  
  
    def __len__(self): return len(self.fpaths)  
  
    def __getitem__(self, ix):  
        f = self.fpaths[ix]  
        target = self.targets[ix]  
        im = (cv2.imread(f)[:,:,:-1])  
        im = cv2.resize(im, (224,224))  
        im = torch.tensor(im/255)  
        im = im.permute(2,0,1)  
        im = self.normalize(im)  
        return im.float().to(device), torch.tensor([target]).float().to(device)
```

```
data = CatsDogs(train_data_dir)
```

Manejo de las partes VGG16

- Fijar la extracción de características: CNN
- Adaptar el clasificador: MLP

```
model = models.vgg16(weights='DEFAULT')
for param in model.parameters():
    param.requires_grad = False
model.avgpool = nn.AdaptiveAvgPool2d(output_size=(1,1))
model.classifier = nn.Sequential(nn.Flatten(),
                                  nn.Linear(512, 128),
                                  nn.ReLU(),
                                  nn.Dropout(0.2),
                                  nn.Linear(128, 1),
                                  nn.Sigmoid())

loss_fn = nn.BCELoss()
optimizer = torch.optim.Adam(model.parameters(), lr= 1e-3)
```


Parámetros de VGG16


Layer (type:depth-idx)	Output Shape	Param #
Sequential: 1-1	[-1, 512, 7, 7]	--
└─Conv2d: 2-1	[-1, 64, 224, 224]	(1,792)
└─ReLU: 2-2	[-1, 64, 224, 224]	--
└─Conv2d: 2-3	[-1, 64, 224, 224]	(36,928)
└─ReLU: 2-4	[-1, 64, 224, 224]	--
└─MaxPool2d: 2-5	[-1, 64, 112, 112]	--
└─Conv2d: 2-6	[-1, 128, 112, 112]	(73,856)
└─ReLU: 2-7	[-1, 128, 112, 112]	--
└─Conv2d: 2-8	[-1, 128, 112, 112]	(147,584)
└─ReLU: 2-9	[-1, 128, 112, 112]	--
└─MaxPool2d: 2-10	[-1, 128, 56, 56]	--
└─Conv2d: 2-11	[-1, 256, 56, 56]	(295,168)
└─ReLU: 2-12	[-1, 256, 56, 56]	--
└─Conv2d: 2-13	[-1, 256, 56, 56]	(590,080)
└─ReLU: 2-14	[-1, 256, 56, 56]	--
└─Conv2d: 2-15	[-1, 256, 56, 56]	(590,080)
└─ReLU: 2-16	[-1, 256, 56, 56]	--
└─MaxPool2d: 2-17	[-1, 256, 28, 28]	--
└─Conv2d: 2-18	[-1, 512, 28, 28]	(1,180,160)
└─ReLU: 2-19	[-1, 512, 28, 28]	--
└─Conv2d: 2-20	[-1, 512, 28, 28]	(2,359,808)
└─ReLU: 2-21	[-1, 512, 28, 28]	--
└─Conv2d: 2-22	[-1, 512, 28, 28]	(2,359,808)
└─ReLU: 2-23	[-1, 512, 28, 28]	--
└─MaxPool2d: 2-24	[-1, 512, 14, 14]	--
└─Conv2d: 2-25	[-1, 512, 14, 14]	(2,359,808)
└─ReLU: 2-26	[-1, 512, 14, 14]	--
└─Conv2d: 2-27	[-1, 512, 14, 14]	(2,359,808)
└─ReLU: 2-28	[-1, 512, 14, 14]	--
└─Conv2d: 2-29	[-1, 512, 14, 14]	(2,359,808)
└─ReLU: 2-30	[-1, 512, 14, 14]	--
└─MaxPool2d: 2-31	[-1, 512, 7, 7]	--
AdaptiveAvgPool2d: 1-2	[-1, 512, 1, 1]	--
Sequential: 1-3	[-1, 1]	--
└─Flatten: 2-32	[-1, 512]	--
└─Linear: 2-33	[-1, 128]	65,664
└─ReLU: 2-34	[-1, 128]	--
└─Dropout: 2-35	[-1, 128]	--
└─Linear: 2-36	[-1, 1]	129
└─Sigmoid: 2-37	[-1, 1]	--
Total params: 14,780,481		
Trainable params: 65,793		
Non-trainable params: 14,714,688		
Total mult-adds (G): 15.36		
Input size (MB): 0.57		
Forward/backward pass size (MB): 103.36		
Params size (MB): 56.38		
Estimated Total Size (MB): 160.32		

Total params: 14,780,481
Trainable params: 65,793
Non-trainable params: 14,714,688
Total mult-adds (G): 15.36

Motivación de la red ResNet

- En la fase hacia adelante en Deep Learning:
 - En las últimas capas se **pierde** casi toda la **información** original
- En la fase de retropropagación (aprendizaje):
 - En las primeras capas, el efecto de la **evanescencia** del **gradiente** es muy acusado
- Solución: añadir una conexión directa, que puentee la entrada con la salida de cada capa (**CONEXIÓN RESIDUAL**)

Conexión Residual

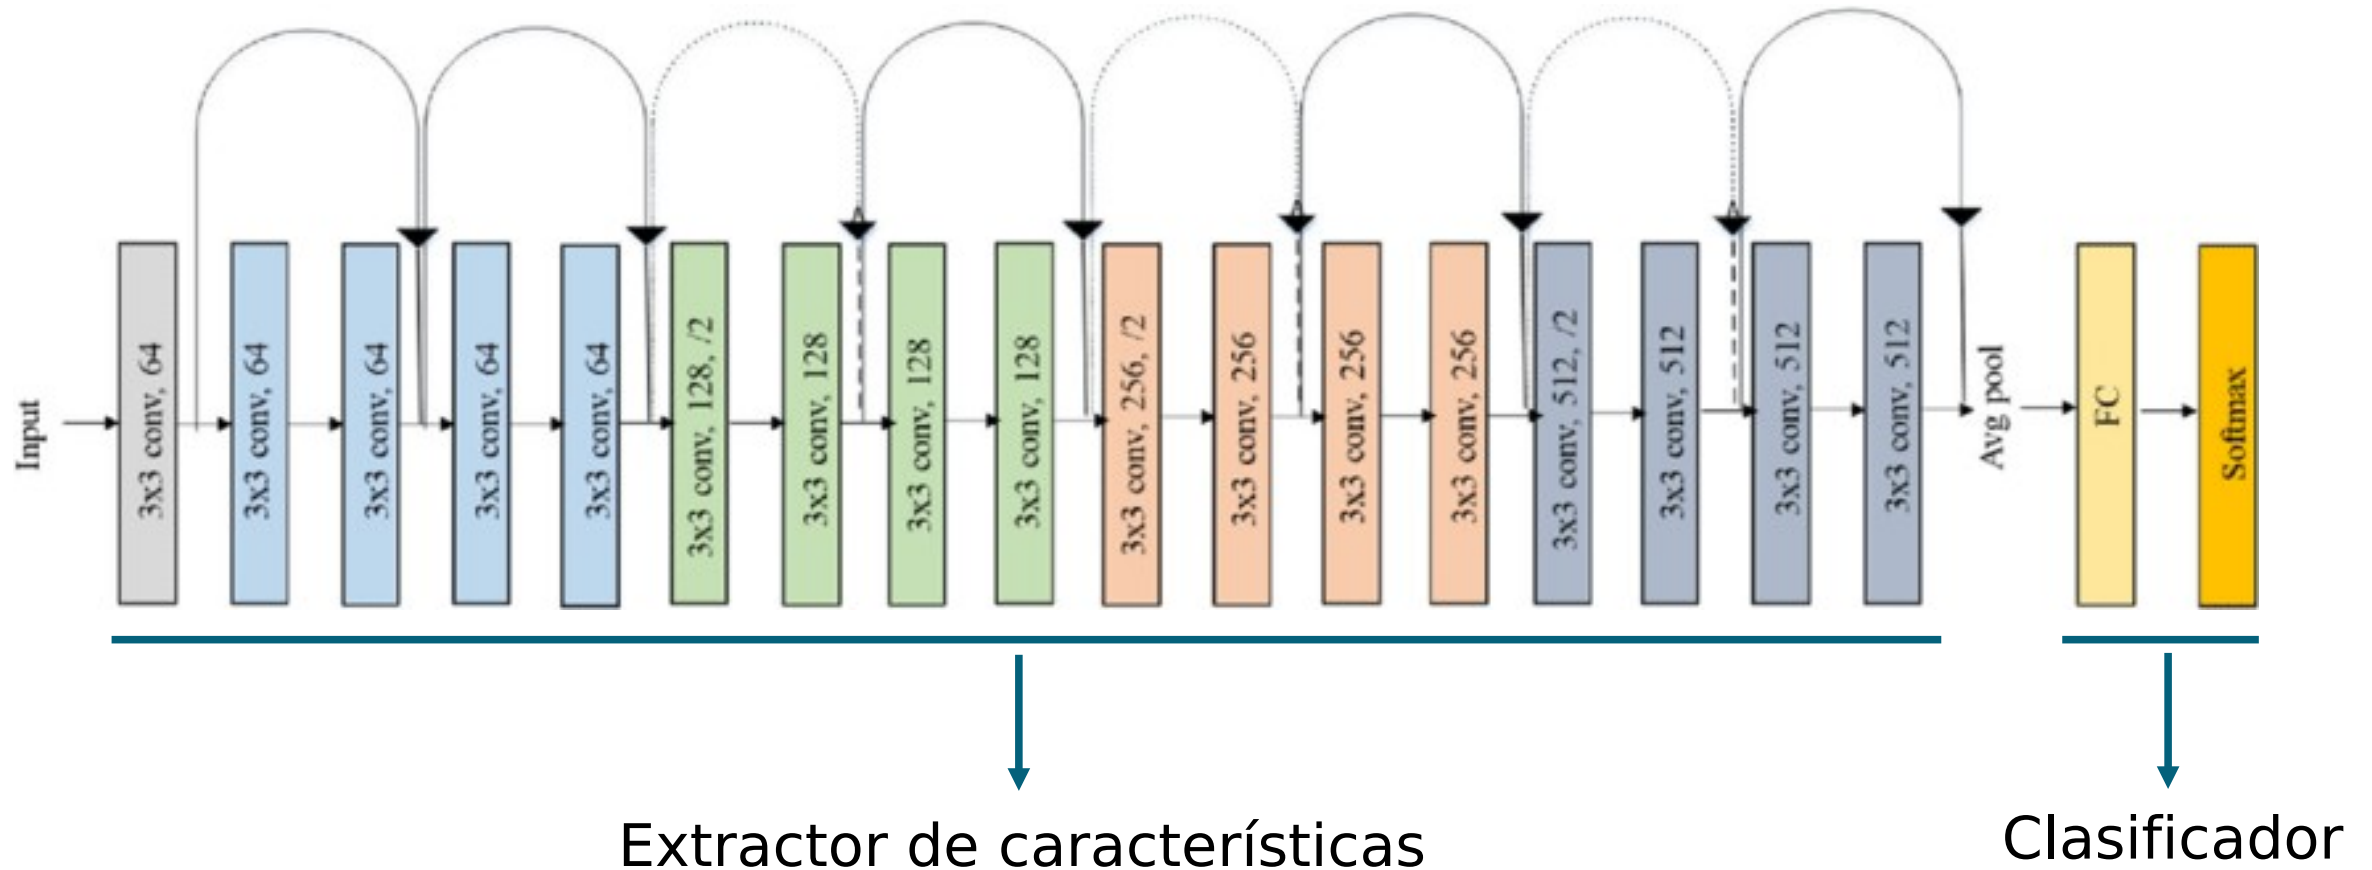


```
class ResLayer(nn.Module):
```

```
    def __init__(self, ni, no, kernel_size, stride=1):
        super(ResLayer, self).__init__()
        padding = kernel_size - 2
        self.conv = nn.Sequential(
            nn.Conv2d(ni, no, kernel_size, stride, padding=padding),
            nn.ReLU())
```

```
    def forward(self, x):
        x = self.conv(x) + x
        return x
```

Arquitectura ResNet



Tipos de VGG y ResNet en Pytorch

- <https://pytorch.org/vision/0.8/models.html?highlight=vgg>
- Simulación:
 - VGG16
 - VGG19
 - ResNet18